

Project 3 - Design Report (Team 28)

Project 3 is a 7-page document.

Frame Table

Basic Descriptions

Limitations and Necessity

Design Proposal

Supplemental Page Table

The CPU is a 32-bit processor. The memory is divided into pages. The Supplemental Page Table (SPT) is a table that maps virtual addresses to physical addresses. It is used to store the physical addresses of the pages that are currently in memory. The SPT is a table of pointers to the physical pages. The SPT is used to store the physical addresses of the pages that are currently in memory. The SPT is a table of pointers to the physical pages. The SPT is used to store the physical addresses of the pages that are currently in memory.

Basic Descriptions

The SPT is a table of pointers to the physical pages. The SPT is used to store the physical addresses of the pages that are currently in memory. The SPT is a table of pointers to the physical pages. The SPT is used to store the physical addresses of the pages that are currently in memory. The SPT is a table of pointers to the physical pages. The SPT is used to store the physical addresses of the pages that are currently in memory.

The SPT is a table of pointers to the physical pages. The SPT is used to store the physical addresses of the pages that are currently in memory. The SPT is a table of pointers to the physical pages. The SPT is used to store the physical addresses of the pages that are currently in memory. The SPT is a table of pointers to the physical pages. The SPT is used to store the physical addresses of the pages that are currently in memory.

The SPT is a table of pointers to the physical pages. The SPT is used to store the physical addresses of the pages that are currently in memory. The SPT is a table of pointers to the physical pages. The SPT is used to store the physical addresses of the pages that are currently in memory. The SPT is a table of pointers to the physical pages. The SPT is used to store the physical addresses of the pages that are currently in memory.

```
/* From devices/block.c. */
/* Reads sector SECTOR from BLOCK into BUFFER, which must
   have room for BLOCK_SECTOR_SIZE bytes.
   Internally synchronizes accesses to block devices, so external
   per-block device locking is unneeded. */
void
block_read (struct block *block, block_sector_t sector, void *buffer)
{
  check_sector (block, sector);
  block->ops->read (block->aux, sector, buffer);
  block->read_cnt++;
}
```

Limitations and Necessity

[illegible]

[illegible]

0000 0000 00000000 0000 0000 00 00 000000 00 00000 00. 00 00000 00 0(key)00 0(value)0000 0000 0000
 0 00 00000 000000, 00 000000 0000 0000 00 00 00 0000 00 00 00000 00 0 0000. 00 00000 00 00 00 00 0000 0
 00 00 00 00 0000, 0 00 00 0000 0000 0000 00 00000. 00 00 00 0000 00000 0000 0000 0000 00 0000 00 00000 0
 0 00000 0000 00. 0000, 00 00000 0000 0000 00 0000 00 0000 0000 0000 0000 00 000000 00 00(hash
 collision)0 00000 0000. 0000, 00000, 0000 Pintos00 000000 00 0000 0000 0 00 00000 0000 00000 0000 00 00000 0
 0000 00 0000 00000 00. 00 00 00 00 000000 0 0000 000000, 0000 000000 00 0000 0000 00 00 000000 00000 0000
 0000 0000.

```
/* Maybe into vm/spt.h. */
struct spte
{
    int size;
    bool swapped;
    void *uaddr;
    mapid_t mapid;
    block_sector_t index;
    struct hash_elem elem;
}
```

3 / 13

00, 00 SPTE 00 0000 00 00 0000 000000, 00 SPTE 000 0000 000(map identifier) 0000 0
 0 mapid 0000 000. 00 00 SPTE 0000 0000 000 00 0000 0000 0000 0000 mapid MAPID_ERROR 0
 0000 00 0000 000 00 0000 0000 0000 0000 0000. 0, 00 0000 00 0000 00 0000 0000000 0000 00 0000
 index 0000 000, 00 0000 00 0000 0000 00 0000 elem 0000 00.

struct process

```

/* From userprog/process.h. */
/* An user process. */
struct process
{
    ...
    /* Owned by vm/spt.c. */
    struct hash spt;
};
  
```

000 process 0000 0 000000 000 0000000 000000, 000 000000 000 0000 SPT 0000 0000 00, SPT 0 000
 0000000 0000 00000000 0 struct process 000 SPT 0000 0 000. 00, 000000 000 00000 SPT 00
 hash_init() 000 00 0000000 00.

page_fault()

```

/* From userprog/exception.c. */
static void
page_fault (struct intr_frame *f)
{
    ...

    /* If it is caused by invalid address passed to the kernel, kill user
    process
        while making no harm to the kernel. */
    if (is_user_vaddr (fault_addr) && !user)
        process_exit (-1);

    /*
        1. Check if current process's SPT has an entry about the faulting
        page.
        2. If the SPT does not have such SPTE, terminate the malicious user
        process.
        3. Else, according to the information in the entry, allocate a
        physical
            frame which the page would be loaded into. This is done with the
        frame
            table.
        - If the physical page frame is full, evict a frame by LRU-
        approximating
            page replacement algorithm.
        - Also, find the swap slot by the swap table.
        - Then, modify the evicted page's SPTE so that it now holds the
  
```

```

block and
    sector where the page is evicted.
4. Load the faulting page into the frame using block_read().
5. Go back to the normal execution flow.
*/
...
}

```

When a page fault occurs, the kernel must find a free frame to load the page into. If there are no free frames, it must evict a page from memory. The kernel must also update the page table to reflect the new state of the page. Finally, the kernel must return control to the user process.

Lazy Loading and Paging

In Pintos, the kernel uses lazy loading to load pages into memory. When a page fault occurs, the kernel only loads the page if it is not already in memory. This is done by calling `block_read()` to load the page from the disk. If the page is already in memory, the kernel simply updates the page table to reflect the new state of the page.

When a page is loaded into memory, it is marked as "dirty" if it has been modified. If a page is dirty, the kernel must write it back to the disk before it can be evicted. This is done by calling `block_write()` to write the page to the disk. If the page is not dirty, the kernel can simply evict it without writing it back to the disk. This is done by calling `block_evict()` to evict the page from memory.

basic descriptions

The kernel uses a simple paging scheme to manage memory. Each page is represented by a page table entry (PTE) in the kernel's page table. The PTE contains the physical address of the page in memory. The kernel uses the PTE to find the page in memory when a page fault occurs.

The kernel also uses a simple paging scheme to manage the disk. Each disk block is represented by a block table entry (BTE) in the kernel's block table. The BTE contains the logical address of the block on the disk. The kernel uses the BTE to find the block on the disk when a page fault occurs.

The kernel uses a simple paging scheme to manage the user process's memory. Each user process has a page table that maps its virtual addresses to physical addresses in memory. The kernel uses the page table to find the page in memory when a page fault occurs.

limitations and necessity

The kernel has several limitations when it comes to memory management. First, it only supports a single level of paging. Second, it only supports a single type of paging (x86-64). Third, it only supports a single type of disk (IDE).

```

/* From userprog/process.c. */
/* Loads an ELF executable from FILE_NAME into the current thread.
   Stores the executable's entry point into *EIP
   and its initial stack pointer into *ESP.
   Returns true if successful, false otherwise. */

```

```

bool
load (const char *file_name, void (**eip) (void), void **esp)
{
    /* Read program headers. */
    file_ofs = ehdr.e_phoff;
    for (i = 0; i < ehdr.e_phnum; i++)
    {
        struct Elf32_Phdr phdr;
        ...
        switch (phdr.p_type)
        {
            ...
            case PT_LOAD:
                ...
                if (!load_segment (file, file_page, (void *) mem_page,
                                   read_bytes, zero_bytes, writable))
                    goto done;
                ...
            }
        }
        ...
    }
}

```

userprog/process.c load() 함수 호출, ELF 파일의 프로그램 헤더를 읽고, load_segment() 함수를 호출하여 각 프로그램 세그먼트를 메모리에 로드합니다. load_segment() 함수는 세그먼트의 시작 위치, 로드할 바이트 수, 그리고 세그먼트가 읽기 전용인지 여부를 인자로 받습니다. load_segment() 함수는 세그먼트를 성공적으로 로드하면 true를 반환하고, 그렇지 않으면 false를 반환합니다.

```

/* From userprog/process.c. */
/* Loads a segment starting at offset ofs in file at address
   upage. In total, read_bytes + zero_bytes bytes of virtual
   memory are initialized, as follows:

   - Read_bytes bytes at upage must be read from file
     starting at offset ofs.

   - Zero_bytes bytes at upage + read_bytes must be zeroed.

   The pages initialized by this function must be writable by the
   user process if writable is true, read-only otherwise.

   Return true if successful, false if a memory allocation error
   or disk read error occurs. */
static bool
load_segment (struct file *file, off_t ofs, uint8_t *upage,
              uint32_t read_bytes, uint32_t zero_bytes, bool writable)
{
    ...
    file_seek (file, ofs);
    while (read_bytes > 0 || zero_bytes > 0)
    {
        /* Calculate how to fill this page.

```

```

        we will read page_read_bytes bytes from file
        and zero the final page_zero_bytes bytes. */
size_t page_read_bytes = read_bytes < pgsize ? read_bytes : pgsize;
size_t page_zero_bytes = pgsize - page_read_bytes;

/* Get a page of memory. */
uint8_t *kpage = palloc_get_page (pal_user);
if (kpage == null)
    return false;

/* Load this page. */
if (file_read (file, kpage, page_read_bytes) != (int)
page_read_bytes)
{
    palloc_free_page (kpage);
    return false;
}
memset (kpage + page_read_bytes, 0, page_zero_bytes);

/* Add the page to the process's address space. */
if (!install_page (upage, kpage, writable))
{
    palloc_free_page (kpage);
    return false;
}

/* Advance. */
read_bytes -= page_read_bytes;
zero_bytes -= page_zero_bytes;
upage += pgsize;
}
return true;
}

```

`load_segment()` 讀取檔案內容到記憶體。它接收 `file` 檔案句柄、`ofs` 偏移量、`read_bytes` 要讀取的字节數。它會調用 `install_page()` 將讀取的數據安裝到用戶空間的 `upage` 地址。如果安裝失敗，它會返回 `false`。如果成功，它會返回 `true`。在 `load_segment()` 中，我們需要讀取 `read_bytes` 字節，並將剩餘的 `zero_bytes` 字節（即 `pgsize - read_bytes`）填充為 0。這確保了整個頁面的數據正確。在 `Pintos` 中，我們需要確保讀取操作不會導致數據損壞，因此我們使用 `write-back` 策略來管理緩存。

在 `userprog/exception.c` 中，我們需要處理用戶空間的異常。在 `Pintos` 中，我們需要確保用戶空間的數據不會被系統代碼修改。因此，我們需要禁用用戶空間的寫入緩存（write-back），並啟用寫回（write-back）策略。這確保了用戶空間的數據在系統代碼執行期間不會被修改。

Design Proposal

在設計中，我們需要確保用戶空間的數據在系統代碼執行期間不會被修改。因此，我們需要禁用用戶空間的寫入緩存（write-back），並啟用寫回（write-back）策略。這確保了用戶空間的數據在系統代碼執行期間不會被修改。在 `load_segment()` 中，我們需要確保讀取操作不會導致數據損壞，因此我們使用 `write-back` 策略來管理緩存。

이 함수는 주어진 파일에서 주어진 오프셋에서 주어진 바이트 수를 읽어와서 주어진 페이지에 복사합니다. 이 함수는 `load_segment()` 함수를 호출하여 주어진 파일에서 주어진 오프셋에서 주어진 바이트 수를 읽어와서 주어진 페이지에 복사합니다.

```
/* From userprog/process.c. */
static bool
load_segment (struct file *file, off_t ofs, uint8_t *upage,
              uint32_t read_bytes, uint32_t zero_bytes, bool writable)
{
    ...
    file_seek (file, ofs);
    while (read_bytes > 0 || zero_bytes > 0)
    {
        /* Calculate how to fill this page.
           we will read page_read_bytes bytes from file
           and zero the final page_zero_bytes bytes. */
        size_t page_read_bytes = read_bytes < PGSIZE ? read_bytes : PGSIZE;
        size_t page_zero_bytes = PGSIZE - page_read_bytes;

        /*
           After calculating the number of bytes to be read and the number of
           bytes
           to be filled with zeros, we should do the following.

           1. Gets the sector number of the underlying block of the file.
           2. Create new SPTE which holds the sector number from the first
           step.
           3. Insert the SPTE with the hash key of upage into the SPT of this
           newly
           executed process.
           4. After this step, the executable file will be demand-paged
           whenever
           the process tries to executed yet unloaded portion of the
           executable,
           by the modified page fault handler discussed above.
        */

        /* Advance. */
        read_bytes -= page_read_bytes;
        zero_bytes -= page_zero_bytes;
        upage += PGSIZE;
    }
    return true;
}
```

이 함수는 주어진 파일에서 주어진 오프셋에서 주어진 바이트 수를 읽어와서 주어진 페이지에 복사합니다. 이 함수는 `load_segment()` 함수를 호출하여 주어진 파일에서 주어진 오프셋에서 주어진 바이트 수를 읽어와서 주어진 페이지에 복사합니다. 이 함수는 `SPT`를 업데이트합니다.

Stack Growth

Basic Descriptions & Limitations

Necessity

Design Proposal

File Memory Mapping

Basic Descriptions

`mmap()` 可將檔案或裝置，透過(file descriptor) `fd` 來映射到記憶體。地址由 `addr` 指定。返回值是映射的起始地址，类型为 `mapped_t`。在 Unix 中，`mmap()` 返回的地址通常是页对齐的。而在 Pintos 中，`mmap()` 返回的地址不一定是页对齐的，这取决于底层的操作系统。在 Pintos 中，`mmap()` 返回的地址总是页对齐的。

`munmap()` 是用于释放之前通过 `mmap()` 分配的内存。在 `munmap()` 调用后，分配的内存区域将不再有效，任何访问该区域的尝试都可能导致未定义行为。在 `SPT` 中，我们使用 `munmap()` 来释放不再需要的内存，以保持内存管理的正确性。

Limitations and Necessity

```
/* From userprog/syscall.c. */
static void
syscall_handler (struct intr_frame *f)
{
  int syscall_number = (int) dereference (f->esp, 0, WORD_SIZE);

  switch (syscall_number) {
    case SYS_HALT: halt (); break;
    case SYS_EXIT: exit (f->esp); break;
    case SYS_EXEC: f->eax = exec (f->esp); break;
    case SYS_WAIT: f->eax = wait (f->esp); break;
    case SYS_CREATE: f->eax = create (f->esp); break;
    case SYS_REMOVE: f->eax = remove (f->esp); break;
    case SYS_OPEN: f->eax = open (f->esp); break;
    case SYS_FILESIZE: f->eax = filesize (f->esp); break;
    case SYS_READ: f->eax = read (f->esp); break;
    case SYS_WRITE: f->eax = write (f->esp); break;
    case SYS_SEEK: seek (f->esp); break;
    case SYS_TELL: f->eax = tell (f->esp); break;
    case SYS_CLOSE: close (f->esp); break;
    /* There's no handling routine for mmap() and munmap()! */
  }
}
```

在 `Pintos` 中，`mmap()` 和 `munmap()` 用于管理虚拟内存。由于 `Pintos` 是一个简单的操作系统，它不支持复杂的内存管理功能，因此我们使用 `mmap()` 和 `munmap()` 来管理进程的虚拟内存空间。

Design Proposal

`syscall_handler()`

```
/* From userprog/syscall.c. */
static void
syscall_handler (struct intr_frame *f)
{
  int syscall_number = (int) dereference (f->esp, 0, WORD_SIZE);

  switch (syscall_number) {
    ...
    case SYS_MMAP: f->eax = mmap (f->esp); break;
    case SYS_MUNMAP: munmap (f->esp); break;
  }
}
```

```

    }
}

```

즉, 이 함수는 사용자 프로그램의 스택에서 1개의 워드 크기의 데이터를 읽어와, 이를 사용자 프로그램의 스택에 매핑합니다.

mmap()

```

/* Maybe into userprog/syscall.c. */
static uint32_t
mmap (void *esp)
{
    int fd = (int) dereference (esp, 1, WORD_SIZE);
    void *addr = (void *) dereference (esp, 2, WORD_SIZE);
    struct file *fp = retrieve_fp (fd);

    if (fp == NULL)
        return MAPID_ERROR;

    /*
     * To implement mmap() system call, this function should do followings;
     *
     * 1. Allocate a mapid by allocate_mapid() call.
     * 2. Get the size of file, divide it by the size of a page (4 KiB). Let
     *    the quotient be N.
     * 3. Add (N + 1) SPTes into the SPT of current process.
     *    - The SPTe should have the underlying block, sector number, and the
     *      size within the page if the size of the file is not aligned with
     *    the size of a page.
     *    - The SPTe should also have the mapid for this mapping.
     * 4. Return the mapid allocated in the first step.
     */
}

```

`mmap()` 함수는 사용자 프로그램의 스택에서 1개의 워드 크기의 데이터를 읽어와, 이를 사용자 프로그램의 스택에 매핑합니다. 이 함수는 15 KiB의 데이터를 읽어와, 이를 4 KiB(= 15 / 4 + 1)의 SPTe의 SPT에 매핑합니다. 이 SPTe의 `size` 필드는 3개의 워드 크기의 데이터를 나타냅니다. 즉, 이 SPTe는 `fd` 필드를 사용하여 사용자 프로그램의 스택에서 데이터를 읽어와, 이를 사용자 프로그램의 스택에 매핑합니다. 이 함수는 사용자 프로그램의 스택에서 데이터를 읽어와, 이를 사용자 프로그램의 스택에 매핑합니다.

`allocate_mapid()`는 사용자 프로그램의 스택에서 데이터를 읽어와, 이를 사용자 프로그램의 스택에 매핑합니다. `threads/thread.c`의 `allocate_tid()`와 `filesys/file.c`의 `allocate_fd()`는 사용자 프로그램의 스택에서 데이터를 읽어와, 이를 사용자 프로그램의 스택에 매핑합니다.

munmap()

```

/* Maybe into userprog/syscall.c. */
static void
munmap (void *esp)

```

`munmap()` 関数は、メモリ領域を解放し、`mapid` 関数で取得した `SPT` のメモリ領域を解放する。このとき、`present` が 0 である場合は、`dirty` が 1 である場合は、`SPT` のメモリ領域を解放する。

Basic Descriptions and Limitations

0 0 00 0000 0000 0000 00 0 000 0000 000 00 swap disk 00. disk 00 0000 0000, 00 0000 physical memory 000 00 000 0 00 0000. 0 0 00000 swap disk 0000 00 swap out, 000 0000 0000 00 swap in 000 00. 00 pintos 000 000 swap disk 0 **struct block** 00 0000 00. 0 block 0 sector(512 B) 000 0000 000, swap-in, out 0 page 000 0000 000 pintos 0000 800 sector 0 00 swap slot 00 00000 000 0 00. 00, 0 000 00 block 0 00 0000 00, **block_get_role()** 0 00 0000 00(**BLOCK_SWAP**) 00 block 0000 000 0 00 00. 00 000 00 00 00 swap disk 0 pintos-mkdisk swap.dsk 0000 --swap-size 000 00 0 00 000 00 swap disk 000 0 00. 000 swap disk 0 000 swap slot 0 page 000 000 0000, 000 000 0 0000 0000 00. swap disk 00 000 physical memory 000 page 0000, 00 0000 000 00000 00 00 0 00 page size(4 KB) 000 swap slot 000 00. 000 swap slot 00, page-aligned 000 00, 000 slot 000 0 0000 0000 00 swap table 000 0 00. 000 swap table 0 slot 0000 000 000 0000 000 00, 00 000 000 000 00(00 00 00) 00 0000 0000 00. 00 pintos 0000 struct block 0 0000 00 000, 00 0000 00 0 000 00 00 0000. 000 000 000 00 00 physical memory 0000 000 0 00 000 0 0000 00 0000 00.

swap table은 swap slot을 관리하는 table이다. swap slot은 swap 공간의 단위이다. swap table은 swap slot을 관리하는 table이다. swap slot은 swap 공간의 단위이다. swap table은 swap slot을 관리하는 table이다. swap slot은 swap 공간의 단위이다.

00. 00, swap disk 00 00 00 executable data 00 000000 0000, 00 0000 slot 0000 000000 0000. 0000 lazy allocation 0000 00 eviction 00(physical memory 0000 0) 000000 slot 0000 000000 00.

0000 swap table 000000 000000 00 swap out, swap in, swap free 000000. 00 swap out 00 memory 0000 page 000000 swap slot 000000 0 0 swap table 0000 000000 swap slot 000000, 00 slot 000000 `block_write()` 0000 page 000000 000000 00, 0000 000000 0000 000000 00 00. 0 0 000000 00 0000 0000 eviction policy 00, 00 physical memory 00 0000 page 0000 00 page 000000 eviction 0 00 0000 00. 0000 0000 page 0000 eviction 00 swap disk 0 000000 0000 0 00. 0000 swap in 0000 000000 000000 page fault 00 0000 swap disk 0 000000 00 page 00 physical memory 00 0000 0000, 0 0 0000 page 000000, 00 page 000000 `block_read()` 0 00 0000 swap slot 000000 000000 00. 00, 00 000000 000000 00 slot 0000 000000 00. 000000 swap free 000000 swap in 000000 000000 physical memory 00 000000 00 00 slot 000000 process 0000 0 0000 000000 00 00 slot 0000 free 0000 00.

00, swap disk 0 000000 00 process 000000 00 000000 000000. 0000 swap table 00 0 000000 000000 000000 0000 00 0000 000000 0 00 overhead 0000 0 00. 0000 000000 vm/swap.c 0 000000 000000 000000, 0 0000 bitmap 0 0 00 disk 0 000000 0000 0 00 0000. 0000 0 00 lock 0 0000 00 bitmap 0 00 0 00 00 0 atomic 0 000000 00 000000 0 0000.

Design Proposal

swap table 00 000000 0000 0000 struct bitmap 000000 00000000 00. vm/swap.c 0 0000 0000 000000 00 00 00 0000 000000 00.

On Process Termination

Basic Descriptions and Limitations

00 0000 000000 00 0000 000000 process 0 0000 0000 0000 000000 00 000000 00 0000 00 000000 00. 0000 00 0000 0000 0000 frame table, supplementary page table, file memory mapping, swap table 00 0000 0000 0000 0 0000 00, lock 0 0000 000000 00 000000 00 00 00 000000 0000 000000 0 000000 000000, 00 000000 userprog/process.c 00 0000 process_exit() 0000 destory_process() 0 00 000000 0000, 00 0000 process 0000 00 file pointer 00 close 0000, directory 00 free 0000 000000 000000 00. 0000 0000 00000000 0000 0000 0000 000000 0 000000 0000, swap slot 00 0000 00 0000 00 000000 00 000000 00. 0000 00 0000 00000000 00.

Necessity and Design Proposal

00 on process termination 0 000000 0000 userprog/process.c 0000, process_exit() 00 0000 0000 0000 0 0. 000000 Supplementary Page Table 000000 00 000000. 000000 00 0000 Page 00000000 0000, 0 00 dirty 0 0000 dirty 0 00 000000 write 00 0 0000 00000000 000000 000000 00. Swap Table 000000 00 process 000000 00 (00 0 0000) swap slot 0000 0000 swap table 0000 00 00000000 000000 00. File Memory Mapping 0 00 000000 close_file() 0 000000, load 00 00 0000 0000 000000 S-page table 00 000000 00 file 0 00 0000 00 000000 00 file 0 close 0000 00 000000 00000000 00. 00, dirty 0 000000 00 000000 00000000 00 write 0000 0000 000000, 0 00 00 memory mapping 00000000 000000 00 process 00 00 file memory mapping 0 00000000 00. 00000000 00 0 00 frame table 00, 0 00 00 00000000 00 process 000000 00 frame 0000 000000 00 000000. 0000 00 process 0 000000 frame 00 frame table 0000 00 000000, 0 0 00 00 frame 00 dirty 0 00000000 (00000000) 00 frame 0 000000 000000 write 00 000000 00000000 00. 00 frame 000000 process 0 0000 0 00 00.